


Rescue: per-turn LMCache hit fraction (n=47 substantive turns)



Quality: LLM-judge A/B on $n = 19$ divergent pairs

32%

32%

37%

■ prefer cacheblend ■ prefer equivalent ■ prefer no_cache